

SMAUG: the Key Exchange Algorithm based on Module-LWE and Module-LWR

Jung Hee Cheon^{1,2}, Hyeongmin Choe¹, Dongyeon Hong⁴, Jeongdae Hong³,
Hyeon Seong², Junbum Shin², and MinJune Yi^{1,2}

¹ Seoul National University
{jhcheon, sixtail528, yiminjune}@snu.ac.kr

² CryptoLab Inc.
{she000, junbum.shin}@cryptolab.co.kr

³ Ministry of National Defense
ghjd2000@gmail.com

⁴ jjoker041@gmail.com

Version 2.0
(October 30, 2023)

Abstract. In this documentation, we introduce a new lattice-based post-quantum key encapsulation mechanism (KEM), a type of key exchange algorithm that we submit to the Korean Post-Quantum Cryptography Competition. Based on the hardness of the MLWE and MLWR problems, defined in module lattices, we design an efficient public key encryption (PKE) scheme **SMAUG.PKE** and KEM scheme **SMAUG.KEM**. With the choice of sparse secret, **SMAUG** achieves ciphertext sizes up to 12% and 9% smaller than Kyber (NIST’s 2022 selection) and Saber (NIST’s finalist), with much faster running time, up to 120% and 70%, respectively.

Keywords: Lattice-based Cryptography · Post-Quantum Cryptography · Key Encapsulation Mechanism · Module Learning With Errors · Module Learning With Roundings.

This work is submitted to ‘Korean Post-Quantum Cryptography Competition’ (www.kpqc.or.kr).

Changelog

October 30, 2023 (version 2.0) First, we update the hamming weight sampler HWT, which was not running in a constant-time due to the dependency on the number of hash calls. The new hamming weight sampler is a hybrid of the previous HWT algorithm, which was adopted from `SampleInBall` algorithm in `Dilithium`, and the constant weight word sampler [53]. We verified that the new sampler runs in a constant time, with a fixed number of hash calls. With the new hamming weight sampler and the partly optimized reference code, `SMAUG` is now 17% faster than previous version.

Second, we give an additionally security analysis for the choice of the approximate discrete Gaussian sampler. Using the Rényi divergence, it is theoretically guaranteed that the security loss comes from the approximation is minute.

Lastly, we give an additionally justification for the decryption failure probability against the state-of-the-art decryption failure attacks, asserting that the current failure probability of `SMAUG` is already low enough due to the attack scenarios.

May 23, 2023 (version 1.0) First, we update the python script for DFP computation as it was computing the decryption failure probability (DFP) wrongly. Note that the script was missing in the submission file, but included in our website. The parameter sets for NIST's security levels 3 and 5 were having higher DFPs then they were reported in the 1st round submission. As a result, the parameter sets are updated.

Second, we additionally compress the ciphertexts. As compression makes the error larger, we exploit the balance between the sizes and DFP.

Third, we put additional cost estimations on some algebraic and topological attacks: Arora-Ge [8], Coded-BKW [38], and Meet-LWE [50] attacks. We note that the previous parameter sets were all in a secure region against these attacks; however, for the new parameter sets, we aim to have more security margins. We put our code for estimating the cost of Meet-LWE attack in the python script.

Based on the above three updates, we changed our recommended parameter sets. As $q = 1024$ is not available anymore for sufficient DFPs in the security levels 3 and 5, we move to $q = 2048$ for those levels, resulting in slightly larger public key and secret key sizes. The ciphertext sizes are decreased by at most 96 bytes.

We also update the reference implementation to have a constant running time with much faster speed. It is uploaded to our website: `kpqc.cryptolab.co.kr/smaug`.

1 Introduction

SMAUG is an efficient post-quantum key encapsulation mechanism whose security is based on the hardness of the lattice problems. The IND-CPA security of SMAUG.PKE relies on the hardness of MLWE problem and MLWR problem, which implies the IND-CCA2 security of SMAUG.KEM.

Our SMAUG.KEM scheme follows the approaches in recent constructions of post-quantum KEMs such as Lizard [26] and RLizard [48]. SMAUG.KEM bases its security on the module variant lattice problems: the public key does not leak the secret key information by the hardness of MLWE problem, and the ciphertext protects sharing keys based on the hardness of MLWR problem. SMAUG consists of an underlying public key encryption (PKE) scheme SMAUG.PKE, which turns into SMAUG.KEM via Fujisaki-Okamoto transform.

1.1 Design rationale

The design rationale of SMAUG aims is to achieve small ciphertext and public key with low computational cost while maintaining security against various attacks. In more detail, we target the following practicality and security requirements considering real applications:

Practicality:

- Both the public key and ciphertext, especially the latter, which is transmitted more frequently, need to be short in order to minimize communication costs.
- As the key exchange protocol is frequently required on various personal devices, a KEM algorithm with low computational costs is more feasible than a high-cost one.
- A small secret key is desirable in restricted environments such as embedded or IoT devices since managing the secure zone is crucial to prevent physical attacks on secret key storage.

Security:

- The shared key should have a large enough entropy, at least ≥ 256 bits, to prevent Grover's search [37].
- Security should be concretely guaranteed concerning the attacks on the underlying assumptions, say lattice attacks.
- The low enough decryption failure probability (DFP) is essential to avoid the attacks boosting the failure and exploiting the decryption failures [28, 44].
- As KEMs are widely used in various devices and systems, countermeasures against implementation-specific attacks should also be considered. Especially combined with DFP, using error correction code (ECC) on the message to reduce decryption failures should be avoided since masking ECC against side-channel attacks is a very challenging problem.

MLWE and MLWR. SMAUG is constructed on the hardness of MLWE (Module-Learning with Errors) and MLWR (Module-Learning with Rounding) problems and follow the key structure of Lizard [26] and Ring-Lizard (RLizard) [48]. Since LWE problem has been a well-studied problem for the last two decades, there are many LWE-based schemes (e.g., Frodo [17]). Ring and module LWE problems are variants defined over structured lattices and regarded as hard as *LWE*. Many schemes base their security on RLWE/MLWE (e.g., NewHope [5], Kyber [16] and Saber [32]) for efficiency reasons. We also chose the module structure, which enables us to fine-tune security and efficiency in a much more scalable way, unlike standard and ring versions. Since MLWR problem is regarded as hard as MLWE problem unless we overuse the same secret to generate the samples [15], we chose to use MLWR samples for the encryption. By basing the security of encryption to MLWR, we reduce the ciphertext size by $\log q / \log p$ than MLWE instances so that more efficient encryption and decryption are possible.

Quantum Fujisaki-Okamoto transform. SMAUG consists of key encapsulation mechanism SMAUG.KEM and public key encryption scheme SMAUG.PKE. On top of the PKE scheme, we construct the KEM scheme using the Fujisaki-Okamoto (FO) transform [35,36]. Line of works on FO transforms in the quantum random oracle model [14,41,45,51] make it possible to analyze the quantum security, i.e., in the quantum random oracle model (QROM). In particular, we use the FO transform with implicit rejection and no ciphertext contributions (FO_m^χ) following [43].

Sparse secret key and ephemeral key. We design the key generation algorithm based on MLWE problem using sparse secret. Based on the hardness reduction on the LWE problem using sparse secret [25], we use sparse ternary polynomials for the secret key and the ephemeral polynomial vectors. We take advantage of the sparsity, e.g., significantly smaller secret keys and faster multiplications. In particular, the small secret makes SMAUG more feasible in IoT devices having restricted resources.

Choice of moduli. All our moduli are powers of 2. This choice makes SMAUG enjoy faster encapsulation using simple bit shiftings, easy uniform sampling, and scalings. The power of 2 moduli makes it hard to apply Number Theoretic Transform (NTT) on the polynomial multiplications. However, small enough moduli and polynomial degrees enable SMAUG to achieve faster speed.

Negligible decapsulation failures. Since we base the security on the lattice problems, noise is inherent. The decryption result of a SMAUG.PKE ciphertext could be different from the original message but with negligible probability, say decryption failure probability (DFP). As other LWE/LWR-based KEMs, SMAUG has the message-independent DFP. We balance the sizes, DFP, and

security of SMAUG by fine-tuning the parameters. Trade-offs between them could give additional parameter sets for specific purposes.

We give estimated security and sizes for our parameter sets in Table 1. The complete parameter sets are given in Section 5. The security is estimated via lattice estimator [2] and our code. It is shown in core-SVP hardness. The DFP is calculated via a python script modified from Lizard and RLizard [42] and is reported in logarithm base two. The sizes are given in bytes. We include the security estimator of SMAUG in the reference code package on the team SMAUG website: kpqc.cryptolab.co.kr/smaug.

Parameters sets Target security	SMAUG 128	SMAUG 192	SMAUG 256
	1	3	5
n	256	256	256
k	2	3	5
(q, p, p')	(1024, 256, 32)	(2048, 256, 256)	(2048, 256, 64)
Classical core-SVP hardness	120.0	181.7	264.5
Quantum core-SVP hardness	105.6	160.9	245.2
Decryption failure probability	-120	-136	-167
Secret key size	176	236	218
Public key size	672	1088	1792
Ciphertext size	672	1024	1472

Table 1: Security and sizes for our parameter sets.

1.2 Advantages and limitations

Advantages The security of SMAUG relies on the hardness of the lattice problems MLWE and MLWR, which enable balancing between security and efficiency. In terms of sizes, SMAUG has smaller ciphertext sizes compared to Kyber or Saber, which is the smallest ciphertext size among the recent practical lattice-based KEMs which avoid using the error correction codes. In terms of DFP, SMAUG achieves low enough DFP, which are similar to that of Saber. We do not use error correction code (ECC) to avoid side-channel attacks targeting decryption failures, to achieve masking-friendly implementation. Implementation-wise, encapsulation and decapsulation of SMAUG can be done efficiently. This makes SMAUG much easier to implement and secure against physical attacks. We give the constant-time C reference code, which proves the completeness and shows the efficiency of SMAUG.

Limitations

- We use MLWR problem, which has been studied shorter than MLWE or LWE problems; however, it has a security reduction to MLWE. MLWE problem with a sparse secret has a similar issue but has been studied much longer and is used in various applications, e.g., homomorphic encryptions.

- As we use MLWE problem for the secret key security, larger public key sizes than Saber are inherent. It can be seen as a trade-off between the public key size versus performance with a smaller secret key size.

2 Preliminaries

2.1 Notation

We denote matrices with bold and upper case letters (e.g., \mathbf{A}) and vectors with bold type and lower case letters (e.g., \mathbf{b}). Unless otherwise stated, the vector is a column vector.

We define a polynomial ring $\mathcal{R} = \mathbb{Z}[x]/(x^n + 1)$ where n is a power of 2 integers and denote a quotient ring by $\mathcal{R}_q = \mathbb{Z}[x]/(q, x^n + 1) = \mathbb{Z}_q[x]/(x^n + 1)$ for a positive integer q . For an integer η , we denote the set of polynomials of degree less than n with coefficients in $[-\eta, \eta] \cap \mathbb{Z}$ as S_η . Let \tilde{S}_η be a set of polynomials of degree less than n with coefficients in $[-\eta, \eta] \cap \mathbb{Z}$. We denote a discrete Gaussian distribution with standard deviation σ as $\mathcal{D}_{\mathbb{Z}, \sigma}$. We define Rényi divergence of order α between to probability distributions P and Q such that $\text{Supp}(P) \subseteq \text{Supp}(Q)$ as $R_\alpha(P\|Q) = \left(\sum_{x \in \text{Supp}(P)} \frac{P(x)^\alpha}{Q(x)^{\alpha-1}} \right)^{1/(\alpha-1)}$, where $\text{Supp}(D)$ for a distribution D is defined as $\text{Supp}(D) = \{x \in D : D(x) \neq 0\}$.

2.2 Lattice assumptions

We first define some well-known lattice assumptions *MLWE* and *MLWR* on the structured Euclidean lattices.

Definition 1 (Decision- $MLWE_{n,q,k,\ell,\eta}$). For positive integers q, k, ℓ, η and the dimension n of \mathcal{R} , we say that the advantage of an adversary \mathcal{A} solving the decision- $MLWE_{n,q,k,\ell,\eta}$ problem is

$$\text{Adv}_{n,q,k,\ell,\eta}^{MLWE}(\mathcal{A}) = \left| \Pr [b = 1 \mid \mathbf{A} \leftarrow \mathcal{R}_q^{k \times \ell}; \mathbf{b} \leftarrow \mathcal{R}_q^k; b \leftarrow \mathcal{A}(\mathbf{A}, \mathbf{b})] \right. \\ \left. - \Pr [b = 1 \mid \mathbf{A} \leftarrow \mathcal{R}_q^{k \times \ell}; (\mathbf{s}, \mathbf{e}) \leftarrow S_\eta^\ell \times S_\eta^k; b \leftarrow \mathcal{A}(\mathbf{A}, \mathbf{A} \cdot \mathbf{s} + \mathbf{e})] \right|$$

Definition 2 (Decision- $MLWR_{n,p,q,k,\ell,\eta}$). For positive integers p, q, k, ℓ, η with $q \geq p \geq 2$ and the dimension n of \mathcal{R} , we say that the advantage of an adversary \mathcal{A} solving the decision- $MLWR_{n,p,q,k,\ell,\eta}$ problem is

$$\text{Adv}_{n,p,q,k,\ell,\eta}^{MLWE}(\mathcal{A}) = \left| \Pr [b = 1 \mid \mathbf{A} \leftarrow \mathcal{R}_p^{k \times \ell}; \mathbf{b} \leftarrow \mathcal{R}_q^k; b \leftarrow \mathcal{A}(\mathbf{A}, \mathbf{b})] \right. \\ \left. - \Pr [b = 1 \mid \mathbf{A} \leftarrow \mathcal{R}_q^{k \times \ell}; \mathbf{s} \leftarrow S_\eta^\ell; b \leftarrow \mathcal{A}(\mathbf{A}, \lfloor p/q \cdot \mathbf{A} \cdot \mathbf{s} \rfloor)] \right|$$

2.3 Public key encryption and key encapsulation mechanism

We then recap the formalisms of PKE and KEM.

Definition 3 (PKE). A public key encryption scheme is a tuple of PPT algorithms $\text{PKE} = (\text{KeyGen}, \text{Enc}, \text{Dec})$ with the following specifications:

- **KeyGen:** a probabilistic algorithm that outputs a public key pk and a secret key sk ;

- **Enc**: a probabilistic algorithm that takes as input a public key pk and a message μ and outputs a ciphertext ct ;
- **Dec**: a deterministic algorithm that takes as input a secret key sk and a ciphertext ct and outputs a message μ .

Let $0 < \delta < 1$. We say that it is $(1 - \delta)$ -correct if for any (pk, sk) generated from KeyGen and μ ,

$$\Pr[\text{Dec}(\text{sk}, \text{Enc}(\text{pk}, \mu)) \neq \mu] \leq \delta,$$

where the probability is taken over the randomness of the encryption algorithm. We call the above probability decryption failure probability (DFP). In addition, we say that it is correct in the (Q)ROM if the probability is taken over the randomness of the (quantum) random oracle, modeling the hash function.

Definition 4 (KEM). A key encapsulation mechanism scheme is a tuple of PPT algorithms $\text{KEM} = (\text{KeyGen}, \text{Encap}, \text{Decap})$ with the following specifications:

- **KeyGen**: a probabilistic algorithm that outputs a public key pk and a secret key sk ;
- **Encap**: a probabilistic algorithm that takes as input a public key pk and outputs a sharing key K and a ciphertext ct ;
- **Decap**: a deterministic algorithm that takes input a secret key sk and a ciphertext ct and outputs a sharing key K .

The correctness of KEM is defined similarly to that of PKE.

We give the advantage function with respect to the attacks against PKE, namely the INDistinguishability under Chosen Plaintext Attacks (IND-CPA).

Definition 5 (IND-CPA security of PKE). For a (quantum) adversary \mathcal{A} against a public key encryption scheme $\text{PKE} = (\text{KeyGen}, \text{Enc}, \text{Dec})$, we define the IND-CPA advantage of \mathcal{A} $= (\mathcal{A}_1, \mathcal{A}_2)$ as follows:

$$\text{Adv}_{\text{PKE}}^{\text{IND-CPA}}(\mathcal{A}) = \left| \Pr_{(\text{pk}, \text{sk})} \left[b = b' \mid \begin{array}{l} (\mu_0, \mu_1, st) \leftarrow \mathcal{A}_1(\text{pk}); \ b \leftarrow \{0, 1\}; \\ \text{ct} \leftarrow \text{Enc}(\text{pk}, \mu_b); \ b' \leftarrow \mathcal{A}_2(\text{pk}, \text{ct}, st) \end{array} \right] - \frac{1}{2} \right|.$$

The probability is taken over the randomness of \mathcal{A} and $(\text{pk}, \text{sk}) \leftarrow \text{KeyGen}(1^\lambda)$.

We then define two advantage functions with respect to the attacks against KEM, namely the INDistinguishability under Chosen Plaintext Attacks (IND-CPA) as in PKE and the INDistinguishability under (adaptively) Chosen Ciphertext Attacks (IND-CCA).

Definition 6 (IND-CPA and IND-CCA security of KEM). For a (quantum) adversary \mathcal{A} against a key encapsulation mechanism $\text{KEM} = (\text{KeyGen}, \text{Encap}, \text{Decap})$, we define the IND-CPA advantage of \mathcal{A} as follows:

$$\text{Adv}_{\text{KEM}}^{\text{IND-CPA}}(\mathcal{A}) = \left| \Pr_{(\text{pk}, \text{sk})} \left[b = b' \mid \begin{array}{l} b \leftarrow \{0, 1\}; \ (K_0, \text{ct}) \leftarrow \text{Encap}(\text{pk}); \\ K_1 \leftarrow \mathcal{K}; \ b' \leftarrow \mathcal{A}(\text{pk}, \text{ct}, K_b) \end{array} \right] - \frac{1}{2} \right|.$$

The probability is taken over the randomness of \mathcal{A} and $(\text{pk}, \text{sk}) \leftarrow \text{KeyGen}(1^\lambda)$. The IND-CCA advantage of \mathcal{A} is defined similarly except that the adversary can query $\text{Decap}(\text{sk}, \cdot)$ oracle on any ciphertext $\text{ct}' (\neq \text{ct})$.

We can then define the (quantum) security notions of PKE and KEM in the (Q)ROM.

Definition 7 ((Q)ROM security of PKE and KEM). *For $T, \epsilon > 0$, we say that a scheme $\mathcal{S} \in \{\text{PKE}, \text{KEM}\}$ is (T, ϵ) -ATK secure in the (Q)ROM if for any (quantum) adversary \mathcal{A} with runtime $\leq T$ given classical access to \mathcal{O} and (quantum) access to a random oracle H , it holds that $\text{Adv}_{\mathcal{S}}^{\text{ATK}}(\mathcal{A}) < \epsilon$, where*

$$\mathcal{O} = \begin{cases} \text{Enc} & \text{if } \mathcal{S} = \text{PKE} \text{ and } \text{ATK} \in \{\text{OW-CPA}, \text{IND-CPA}\}, \\ \text{Encap} & \text{if } \mathcal{S} = \text{KEM} \text{ and } \text{ATK} = \text{IND-CPA}, \\ \text{Encap}, \text{Decap}(\text{sk}, \cdot) & \text{if } \mathcal{S} = \text{KEM} \text{ and } \text{ATK} = \text{IND-CCA}. \end{cases}$$

2.4 Fujisaki-Okamoto transform

Fujisaki and Okamoto proposed a novel generic transform [35, 36] that turns a weakly secure PKE scheme into a strongly secure PKE scheme in the Random Oracle Model (ROM), and various variants have been proposed to deal with tightness, non-correct PKEs, and in the quantum setting, i.e., QROM. Here, we recall the FO transformation for KEM as introduced by Dent [30] and revisited by Hofheinz et al. [41], Bindel et al. [13], and Hövelmanns et al. [43].

The original FO transforms FO_m^\perp constructs a KEM from a deterministic PKE, i.e., a de-randomized version. The encapsulation randomly samples a message m and uses the message’s hash value $G(m)$ as randomness for encryption, generating a ciphertext. The sharing key $K = H(m)$ is generated by hashing (with different hash functions) the message. In the decapsulation, it first decrypts the ciphertext and recovers the message, m' . If it fails to decrypt, it outputs \perp . If the “re-encryption” of the recovered message is not equal to the received ciphertext, it also outputs \perp . The sharing key can be generated by hashing the recovered message.

In the quantum setting, however, the FO transform with “implicit rejection” (FO_m^\perp) has a tighter security proof than the original version, which implicitly outputs a pseudo-random sharing key if the re-encryption fails. We recap the QROM proof of Bindel et al. [13] allowing the KEMs constructed over non-perfect PKEs to have IND-CCA security.

Theorem 1 ([13], Theorem 1 & 2). *Let G and H be quantum-accessible random oracles, and the deterministic PKE is ϵ -injective. Then the advantage of IND-CCA attacker \mathcal{A} with at most Q_{Dec} decryption queries and Q_G and Q_H hash queries at depth at most d_G and d_H , respectively, is*

$$\begin{aligned} \text{Adv}_{\text{KEM}}^{\text{IND-CCA}}(\mathcal{A}) \leq & 2\sqrt{(d_G + 2) \left(\text{Adv}_{\text{PKE}}^{\text{IND-CPA}}(\mathcal{B}_1) + 8(Q_G + 1)/|\mathcal{M}| \right)} \\ & + \text{Adv}_{\text{PKE}}^{\text{DF}}(\mathcal{B}_2) + 4\sqrt{d_H Q_H / |\mathcal{M}|} + \epsilon, \end{aligned}$$

where \mathcal{B}_1 is an IND-CPA adversary on PKE and \mathcal{B}_2 is an adversary against finding a decryption failing ciphertext, returning at most Q_{Dec} ciphertexts.

3 Design choices

In this section, we explain the design choices for SMAUG.

3.1 MLWE public key and MLWR ciphertext

One of the core designs of SMAUG uses the MLWE hardness for its secret key security and MLWR hardness for its message security. This choice is adapted from Lizard and RLizard, which use LWE/LWR and RLWE/RLWR, respectively. Using both LWE and LWR variant problems makes the conceptual security distinction between the secret key and the ephemeral sharing key: a more conservative secret key with more efficient en/decapsulations. This can be viewed as a trade-off between “conservative” and “efficient” designs. Combined with the sparse secret, bringing the LWE-based key generation to the LWR-based scheme enables balancing the speed and the DFP.

3.1.1 Public key. Public key of SMAUG consists of a vector \mathbf{b} over a polynomial ring \mathcal{R}_q and a matrix \mathbf{A} , which can be viewed as an MLWE sample,

$$(\mathbf{A}, \mathbf{b} = -\mathbf{A}^\top \mathbf{s} + \mathbf{e}) \in \mathcal{R}_q^{k \times k} \times \mathcal{R}_q^k,$$

where \mathbf{s} is a ternary secret polynomial with hamming weight h_s , and \mathbf{e} is an error sampled from discrete Gaussian distribution with standard deviation σ . We hereby specify the uniform matrix sampling algorithm for $\mathbf{A} \in \mathcal{R}_q^{k \times k}$ in Figure 1. It is adapted from the pseudorandom generator `gen` in Saber [29].

<code>expandA(seed):</code> 1: <code>buf ← XOF(seed)</code> 2: for i from 0 to $k - 1$ do 3: <code>A[i] = bytes_to_Rq(buf + polybytes · i)</code> 4: return <code>A</code>	$\triangleright \text{seed} \in \{0, 1\}^{256}$ $\triangleright \text{Convert to ring elements}$
---	---

Fig. 1: Uniform random matrix sampler, `expandA`.

We note that the public key of SMAUG consists of \mathbf{b} and the seed of \mathbf{A} .

3.1.2 Ciphertext. The ciphertext of SMAUG is a tuple of a vector $\mathbf{c}_1 \in \mathcal{R}_p^k$ and a polynomial $c_2 \in \mathcal{R}_p$. The ciphertext is generated by multiplying a random vector \mathbf{r} to the public key; then it is scaled and rounded as,

$$\mathbf{c} = \begin{bmatrix} \mathbf{c}_1 \\ c_2 \end{bmatrix} = \left\lfloor \frac{p}{q} \cdot \begin{pmatrix} \mathbf{A} \\ \mathbf{b}^\top \end{pmatrix} \cdot \mathbf{r} \right\rfloor + \frac{p}{t} \cdot \begin{bmatrix} 0 \\ \mu \end{bmatrix},$$

Along with the public key, it can be treated as an MLWR sample added by a scaled message as $(\mathbf{A}', \lfloor p/q \cdot \mathbf{A}' \cdot \mathbf{r} \rfloor) + (0, \mu')$, where \mathbf{A}' is a concatenated matrix of \mathbf{A} and \mathbf{b}^\top .

The ciphertext can be further compressed by scaling the second component c_2 by p'/p , resulting in a shorter ciphertext but a larger error. We note that the public key can be compressed with the same technique. However, it introduces a more significant error, so we do not compress the public key in SMAUG.

3.2 Sparse secret

For the randomnesses \mathbf{s} and \mathbf{r} , we use the sparse ternary distribution. In the following, we will discuss the advantages of the sparse secret and give the sampling algorithm.

3.2.1 Advantage of using sparse secret The sparse secret is widely used in homomorphic encryption to reduce the noise propagation during the homomorphic operations [19, 24, 39] and to speed up the computations. As the lattice-based KEM schemes have inherent decryption error from LWE or LWR noise, the sparse secret can lower this decryption error and also improve the performance of KEMs.

Concretely, the decryption error can be expressed as $\langle \mathbf{e}, \mathbf{r} \rangle + \langle \mathbf{e}_1, \mathbf{s} \rangle + e_2$, where \mathbf{s} is a secret key, \mathbf{r} is a randomness used for encryption, $\mathbf{e} \leftarrow \chi_{pk}^k$ is a noise added in public key, and $(\mathbf{e}_1, e_2) \leftarrow \chi_{ct}^{k+1}$ is a noise added in ciphertext. As the vectors \mathbf{r} and \mathbf{s} are binary (ternary, resp.), each coefficient of the decryption error is an addition (signed addition, resp.) of h_r variables from χ_{pk} and $h_s + 1$ variables from χ_{ct} . The magnitude of the decryption error depends greatly on the Hamming weights h_r and h_s , and thus, we can take advantage of the sparse secrets.

Other major advantages of sparse secrets include reducing the secret key size and enabling fast polynomial multiplication. As the coefficients of the secret key are sparse with a fixed hamming weight, we can store only the information of the nonzero coefficients. We can further use this structure for the polynomial multiplications, which we will describe in Section 3.4.

On the other hand, as the sparse secret reduces the secret key entropy, the hardness of the lattice problem may be decreased. For the security of LWE problem using sparse secret, a series of works have been done, including [25] for asymptotic security based on the reductions to worst-case lattice problems, and [12, 33, 54] for concrete security. Independent of the secret distribution, the module variant (MLWE) is regarded as hard as LWE problem with appropriate parameters, including a smaller modulus. We also exploit the reductions from ordinary MLWE to MLWE using sparse secret or small errors [20]. The MLWR problem also has a simple reduction from MLWE independent of the secret distribution, and its concrete security is heuristically discussed in [29].

Since SMAUG uses a sparse secret key \mathbf{s} and a sparse randomness \mathbf{r} , the security of SMAUG is based on the hardness of MLWE and MLWR problems using

sparse secret. For the specific parameters, we exploit the lattice-estimator [2], which covers most of the recent lattice attacks, and also consider some attacks not included in the estimator. Using a smaller modulus, SMAUG can maintain high security, as in Kyber or Saber.

3.2.2 Hamming weight sampler Our hamming weight sampler, HWT_h , is a hybrid of the `SampleInBall` algorithm in Dilithium [31] and the `CWW` (constant weight word) sampler in BIKE [53], which have a constant running time. A detailed algorithm is given in Figure 2, which samples a ternary polynomial vector having a hamming weight of h .

$\text{HWT}_h(\text{seed}):$ 1: $\text{idx} = 0$ 2: $(\text{buf}, \text{rand}) \leftarrow \text{XOF}(\text{seed})$ 3: for i from $n - h$ to $n - 1$ do 4: $\text{degree} = \lfloor (\text{buf}[\text{idx}] \cdot (i + 1)) / 2^{32} \rfloor$ 5: $\text{res}[i] = \text{res}[\text{degree}]$ 6: $\text{res}[\text{degree}] = (-1)^{\text{rand}[\text{idx}]}$ 7: $\text{idx} = \text{idx} + 1$ 8: return $\text{convToIdx}(\text{res})$	$\triangleright \text{seed} \in \{0, 1\}^{256}$ $\triangleright \text{buf}[\text{idx}] \in \{0, 1\}^{32}$ $\triangleright \text{rand}[\text{idx}] \in \{0, 1\}$ \triangleright Storing the indexes
---	---

Fig. 2: Hamming weight sampler, HWT_h .

3.3 Discrete Gaussian noise

3.3.1 Using approximate discrete Gaussian noise Our design choice for the noise distribution in MLWE follows the conventional discrete Gaussian distribution, but with approximated CDTs following the approaches in FrodoKEM [17]. As a result, we use a discrete Gaussian noise for the public key generation, which is approximated to a narrow distribution. As this approximated discrete Gaussian noise is used only for the public key, we can efficiently bound the security loss. We give a theoretical justification based on Rényi divergence to guarantee the security of SMAUG, considering the narrow discrete Gaussian noise.

In SMAUG, the narrow discrete Gaussian noise is used only for the public key generation. So the difference in the noise distribution only affects the distinguishing advantage between the games G_2 and G_3 in the proof of Theorem 4. Then the bound for the distinguishing advantage can also be expressed as

$$\left(\text{Adv}_{n,q,k,k,\mathcal{D}_{\mathbb{Z},\sigma}}^{\text{MLWE}}(\mathcal{B}_2) \cdot R_\alpha(\text{dGaussian}_\sigma \| \mathcal{D}_{\mathbb{Z},\sigma})^{nk} \right)^{1-1/\alpha},$$

assuming the pseudorandomness of dGaussian_σ . This is due to Lemma 5.5 in [4]. We note that the key generation calls dGaussian only nk times and that the public key is generated only once.

The bound for our typical parameter set (see Section 5.2) for security level 1 (levels 3 and 5, resp.) increases from $2^{-120.0}$ ($2^{-181.7}$ and $2^{-264.5}$, resp.) to $2^{-118.2}$ ($2^{-176.9}$ and $2^{-260.2}$, resp.) with $\alpha = 200$ (75 and 200, resp.). Opposed to the estimated security based on the bound $\text{Adv}_{n,q,k,k,\text{dGaussian}_\sigma}^{\text{MLWE}}(\mathcal{B}_2)$ given in Section 5.2, this new bound provides a more conservative security preventing some possible future attacks that target the noise distribution.

By using one more bit for **dGaussian** algorithm, we can decrease the advantage to $2^{-119.6}$ ($2^{-181.2}$ and $2^{-263.6}$, resp.) with $\alpha = 500$. This modification will slightly decrease only the speed of key generation by less than 1.1x.

We note that the narrow Gaussian noise is already taken into account when estimating the concrete security (given in Section 5.2) using the explained estimators. The analysis given here provides a more conservative security, preventing some possible future attacks that target the noise distribution. We also note that in the core-SVP methodology, we only focus on the estimated attack cost of the underlying MLWE and MLWR problems, not based on the security reductions (as done in most of the NIST-submitted schemes) for a fair comparison to Kyber or Saber.

3.3.2 dGaussian sampler We construct **dGaussian**, a constant-time approximate discrete Gaussian noise sampler, upon a Cumulative Distribution Table (CDT) but is not used during sampling, as it is expressed with bit operations. We first scale the discrete Gaussian distribution and make a CDT approximating the discrete Gaussian distribution. We choose an appropriate scaling factor based on the analysis in [17, 47] using Rényi divergence. We then deploy the Quine-McCluskey method⁵ and apply logic minimization technique on the CDT. As a result, even though our **dGaussian** is constructed upon CDT, it is expressed with bit operations and is constant-time. The algorithms are easily parallelizable and suitable for IoT devices as their memory requirement is low.

We describe **dGaussian** with $\sigma = 1.0625$ in Figure 3 and $\sigma = 1.453713$ in Figure 4.

dGaussian_{1.0625}(x):

Require: $x = x_0x_1x_2x_3x_4x_5x_6x_7x_8x_9 \in \{0, 1\}^{10}$

- 1: $s = s_1s_0 = 00 \in \{0, 1\}^2$
- 2: $s_0 = x_0x_1x_2x_3x_4x_5x_7\overline{x_8}$
- 3: $s_0 += (x_0x_3x_4x_5x_6x_8) + (x_1x_3x_4x_5x_6x_8) + (x_2x_3x_4x_5x_6x_8)$
- 4: $s_0 += (\overline{x_2x_3x_6x_8}) + (\overline{x_1x_3x_6x_8})$
- 5: $s_0 += (x_6x_7\overline{x_8}) + (\overline{x_5x_6x_8}) + (\overline{x_4x_6x_8}) + (\overline{x_7x_8})$
- 6: $s_1 = (x_1x_2x_4x_5x_7x_8) + (x_3x_4x_5x_7x_8) + (x_6x_7x_8)$
- 7: $s = (-1)^{x_9} \cdot s$ $\triangleright \cdot$ is the arithmetic multiplication
- 8: **return** s

Fig. 3: Discrete Gaussian sampler with $\sigma = 1.0625$, **dGaussian** _{σ} .

⁵ We use the python package, from <https://github.com/dreylago/logicmin>.

```

dGaussian1.453713(x):
Require:  $x = x_0x_1x_2x_3x_4x_5x_6x_7x_8x_9x_{10} \in \{0, 1\}^{11}$ 
1:  $s = s_2s_1s_0 = 000 \in \{0, 1\}^3$ 
2:  $s_0 = (x_0x_1x_2x_3x_5x_7x_8) + (x_1x_2x_3x_5\overline{x_6x_7x_9}) + (\overline{x_1x_2x_3x_6x_7x_8})$ 
3:  $s_0 += (\overline{x_1x_2x_3x_5x_8x_9}) + (\overline{x_0x_2x_3x_5x_8x_9})$ 
4:  $s_0 += (x_4x_5\overline{x_6x_7x_9}) + (x_3x_4x_8\overline{x_9}) + (\overline{x_5x_6x_7x_8}) + (\overline{x_4x_6x_7x_8}) + (\overline{x_4x_5x_8x_9})$ 
5:  $s_0 += (x_5x_8\overline{x_9}) + (x_6x_8\overline{x_9}) + (x_7x_8\overline{x_9}) + (\overline{x_7x_8x_9}) + (\overline{x_6x_8x_9})$ 
6:  $s_1 = (x_0x_1x_4\overline{x_5x_6x_7x_9}) + (x_2x_4\overline{x_5x_6x_7x_9}) + (x_3x_4\overline{x_5x_6x_7x_9}) + (x_5x_6x_7\overline{x_8x_9})$ 
7:  $s_1 += (\overline{x_1x_2x_3x_8x_9}) + (\overline{x_7x_8x_9}) + (\overline{x_6x_8x_9}) + (\overline{x_5x_8x_9}) + (\overline{x_4x_8x_9})$ 
8:  $s_2 = (x_1x_4x_5x_6x_7x_8x_9) + (x_2x_4x_5x_6x_7x_8x_9) + (x_3x_4x_5x_6x_7x_8x_9)$ 
9:  $s = (-1)^{x_{10}} \cdot s$   $\triangleright \cdot$  is the arithmetic multiplication
10: return  $s$ 

```

Fig. 4: Discrete Gaussian sampler with $\sigma = 1.453713$, dGaussian _{σ} .

3.4 Polynomial multiplication using sparsity

SMAUG uses the power-of-two moduli to ease the correct scaling and roundings. However, this makes the polynomial multiplications hard to benefit from Number Theoretic Transform (NTT). To address this issue, we propose a new polynomial multiplication that takes advantage of sparsity, which we adapt from [1, 48]. Our new multiplication, given in Figure 5, is constant-time and is faster than the previous approach. We also use a similar secret storing technique as RLizard, where only the degrees of non-zero coefficients are stored in the secret key and directly used in polynomial multiplications.

```

poly_mult_add(a, b, neg_start):  $\triangleright a \in \mathcal{R}_q, b \in \mathcal{S}_\eta$ 
1:  $c = 0$ 
2: for  $i$  from 0 to  $\text{neg\_start} - 1$  do
3:    $\text{degree} = b[i]$ 
4:   for  $j$  from 0 to  $n - 1$  do
5:      $c[\text{degree} + j] = c[\text{degree} + j] + a[j];$ 
6: for  $i$  from  $\text{neg\_start}$  to  $\text{len}(b) - 1$  do
7:    $\text{degree} = b[i]$ 
8:   for  $j$  from 0 to  $n - 1$  do
9:      $c[\text{degree} + j] = c[\text{degree} + j] - a[j];$ 
10: for  $j$  from 0 to  $n - 1$  do
11:    $c[j] = c[j] - c[n + j];$ 
12: return  $c$ 

```

Fig. 5: Polynomial multiplication using sparsity.

3.5 FO transform, FO_m^χ

We construct SMAUG upon the FO transform with implicit rejection and without ciphertext contribution to the sharing key generation, say FO_m^χ . This choice makes the encapsulation and decapsulation algorithm efficient since the sharing key can be directly generated from a message. The public key is additionally fed into the hash function with the message to avoid multi-target decryption failure attacks. The IND-CCA security of the resulting KEM in the QROM is well-studied in [13, 41, 43].

4 The SMAUG

4.1 Specification of SMAUG.PKE

We now describe the public key encryption scheme SMAUG.PKE in Figure 6 with the following building blocks:

- Extendable output function XOF for generating seed_A , seed_{sk} , and seed_e ,
- Uniform random matrix sampler expandA for deriving \mathbf{A} from seed_A ,
- Discrete Gaussian sampler dGaussian_σ for deriving a MLWE noise \mathbf{e} with standard deviation σ from seed_e ,
- Hamming weight sampler HWT_h for deriving a sparse ternary \mathbf{s} (resp. \mathbf{r}) with hamming weight $h = h_s$ (resp. $h = h_r$) from seed_{sk} (resp. seed_r).

KeyGen(1^λ): 1: $\text{seed} \leftarrow \{0, 1\}^{256}$ 2: $(\text{seed}_A, \text{seed}_{sk}, \text{seed}_e) \leftarrow \text{XOF}(\text{seed})$ 3: $\mathbf{A} \leftarrow \text{expandA}(\text{seed}_A) \in \mathcal{R}_q^{k \times k}$ 4: $\mathbf{s} \leftarrow \text{HWT}_{h_s}(\text{seed}_{sk}) \in S_\eta^k$ 5: $\mathbf{e} \leftarrow \text{dGaussian}_\sigma(\text{seed}_e) \in \mathcal{R}^k$ 6: $\mathbf{b} = -\mathbf{A}^\top \cdot \mathbf{s} + \mathbf{e} \in \mathcal{R}_q^k$ 7: return $\text{pk} = (\text{seed}_A, \mathbf{b})$, $\text{sk} = \mathbf{s}$	
Enc($\text{pk}, \mu; \text{seed}_r$): 1: $\mathbf{A} = \text{expandA}(\text{seed}_A)$ 2: if seed_r is not given then $\text{seed}_r \leftarrow \{0, 1\}^{256}$ 3: $\mathbf{r} \leftarrow \text{HWT}_{h_r}(\text{seed}_r) \in S_\eta^k$ 4: $\mathbf{c}_1 = \lfloor p/q \cdot \mathbf{A} \cdot \mathbf{r} \rfloor \in \mathcal{R}_p^k$ 5: $c_2 = \lfloor p'/q \cdot \langle \mathbf{b}, \mathbf{r} \rangle + p'/t \cdot \mu \rfloor \in \mathcal{R}_{p'}$ 6: return $\text{ct} = (\mathbf{c}_1, c_2)$	$\triangleright \text{pk} = (\text{seed}_A, \mathbf{b}), \mu \in \mathcal{R}_t$
Dec(sk, \mathbf{c}): 1: $\mu' = \lfloor t/p \cdot \langle \mathbf{c}_1, \mathbf{s} \rangle + t/p' \cdot c_2 \rfloor \in \mathcal{R}_t$ 2: return μ'	$\triangleright \text{sk} = \mathbf{s}, \mathbf{c} = (\mathbf{c}_1, c_2)$

Fig. 6: Description of SMAUG.PKE

We then prove the completeness of SMAUG.PKE.

Theorem 2 (Completeness of SMAUG.PKE). *Let \mathbf{A} , \mathbf{b} , \mathbf{s} , \mathbf{e} , and \mathbf{r} are defined as in Figure 6. Let the moduli t , p , p' , and q satisfy $t \mid p \mid q$ and $t \mid p' \mid q$. Let $\mathbf{e}_1 \in \mathcal{R}_\mathbb{Q}^k$ and $e_2 \in \mathcal{R}_\mathbb{Q}$ be the rounding errors introduced from the scalings and roundings of $\mathbf{A} \cdot \mathbf{r}$ and $\mathbf{b}^T \cdot \mathbf{r}$. That is, $\mathbf{e}_1 = \frac{q}{p}(\lfloor \frac{p}{q} \cdot \mathbf{A} \cdot \mathbf{r} \rfloor \bmod p) - (\mathbf{A} \cdot \mathbf{r} \bmod q)$ and $e_2 = \frac{q}{p'}(\lfloor \frac{p'}{q} \cdot \langle \mathbf{b}, \mathbf{r} \rangle \rfloor \bmod p') - (\langle \mathbf{b}, \mathbf{r} \rangle \bmod q)$. Let $\delta =$*

$\Pr [\|\langle \mathbf{e}, \mathbf{r} \rangle + \langle \mathbf{e}_1, \mathbf{s} \rangle + e_2\|_\infty > \frac{q}{2t}]$, where the probability is taken over the randomness of the encryption. Then **SMAUG.PKE** in Figure 6 is $(1 - \delta)$ -correct. That is, for every message μ and every key-pair $(\mathbf{pk}, \mathbf{sk})$ returned by $\text{KeyGen}(1^\lambda)$, the decryption fails with a probability less than δ .

Proof. By the definition of \mathbf{e}_1 and e_2 , it holds that $\mathbf{c}_1 = \frac{p}{q} \cdot (\mathbf{A} \cdot \mathbf{r} + \mathbf{e}_1) \pmod{p}$ and $c_2 = \frac{p'}{q} \cdot (\langle \mathbf{b}, \mathbf{r} \rangle + e_2) + \frac{p'}{t} \cdot \mu \pmod{p'}$, where the coefficients of \mathbf{e}_1 and e_2 are in $\mathbb{Z} \cap (-\frac{q}{2p}, \frac{q}{2p}]$ and $\mathbb{Z} \cap (-\frac{q}{2p'}, \frac{q}{2p'}]$, respectively. Thus, the decryption of the ciphertext (\mathbf{c}_1, c_2) can be written as

$$\begin{aligned} \left\lfloor \frac{t}{p} \cdot \langle \mathbf{c}_1, \mathbf{s} \rangle + \frac{t}{p'} \cdot c_2 \right\rfloor \pmod{t} &= \left\lfloor \frac{t}{q} (\langle \mathbf{A} \cdot \mathbf{r}, \mathbf{s} \rangle + \langle \mathbf{e}_1, \mathbf{s} \rangle + \langle \mathbf{b}, \mathbf{r} \rangle + e_2) + \mu \right\rfloor \pmod{t} \\ &= \left\lfloor \frac{t}{q} (\langle \mathbf{A}^\top \cdot \mathbf{s} + \mathbf{b}, \mathbf{r} \rangle + \langle \mathbf{e}_1, \mathbf{s} \rangle + e_2) + \mu \right\rfloor \pmod{t} \\ &= \mu + \left\lfloor \frac{t}{q} (\langle \mathbf{e}, \mathbf{r} \rangle + \langle \mathbf{e}_1, \mathbf{s} \rangle + e_2) \right\rfloor \pmod{t}. \end{aligned}$$

This is equal to μ if and only if every coefficient of $\langle \mathbf{e}, \mathbf{r} \rangle + \langle \mathbf{e}_1, \mathbf{s} \rangle + e_2$ is in the interval $[-\frac{q}{2t}, \frac{q}{2t})$. It concludes the proof. \square

4.2 Specification of SMAUG.KEM

We finally introduce the key encapsulation mechanism **SMAUG.KEM** in Figure 7. **SMAUG.KEM** is designed following the Fujisaki-Okamoto transform with implicit rejection using the non-perfectly correct public key encryption **SMAUG.PKE**. The construction of **SMAUG.KEM** involves the use of the following symmetric primitives:

- Hash function H for hashing a public key,
- Hash function G for deriving a sharing key and a seed.

The Fujisaki-Okamoto transform used in Figure 7 defers from the FO_m^χ transform in [43] in encapsulation and decapsulation. When generating the sharing key and randomness, **SMAUG's Encap** utilizes the hashed public key, which prevents certain multi-target attacks. As for **Decap**, if $\text{ct} \neq \text{ct}'$ holds, an alternative sharing key should be re-generated not to leak failure information against Side-Channel Attacks (SCA). However, even when the failure information is leaked, security can still rely on the explicit FO transform FO_m^\perp , recently treated in [41] with a competitive bound.

We also remark that the randomly chosen message μ should be hashed additionally in the environments using a non-cryptographic system Random Number Generator (RNG). Using a True Random Number Generator (TRNG) is recommended to sample the message μ in such devices.

We now show the completeness of **SMAUG.KEM** based on the completeness of the underlying public key encryption scheme, **SMAUG.PKE**.

KeyGen(1^λ): 1: $(pk, sk') \leftarrow \text{SMAUG.PKE.KeyGen}(1^\lambda)$ 2: $d \leftarrow \{0, 1\}^{256}$ 3: return $pk, sk = (sk', d)$	
Encap(pk): 1: $\mu \leftarrow \{0, 1\}^{256}$ 2: $(K, \text{seed}) \leftarrow G(\mu, H(pk))$ 3: $ct \leftarrow \text{SMAUG.PKE.Enc}(pk, \mu; \text{seed})$ 4: return ct, K	$\triangleright pk = (\text{seed}_A, b)$
Decap(sk, ct): 1: $\mu' = \text{SMAUG.PKE.Dec}(sk', ct)$ 2: $(K', \text{seed}') \leftarrow G(\mu', H(pk))$ 3: $ct' = \text{SMAUG.PKE.Enc}(pk, \mu'; \text{seed}')$ 4: if $ct \neq ct'$ then 5: $(K', \cdot) \leftarrow G(d, H(ct))$ 6: return K'	$\triangleright sk = (sk', d)$

Fig. 7: Description of SMAUG.KEM

Theorem 3 (Completeness of SMAUG.KEM). *We borrow the notations and assumptions from Theorem 2 and Figure 7. Then SMAUG.KEM in Figure 7 is also $(1 - \delta)$ -correct. That is, for every key-pair (pk, sk) generated by $\text{KeyGen}(1^\lambda)$, the shared keys K and K' are identical with probability larger than $1 - \delta$.*

Proof. The shared keys K and K' are identical if the decryption succeeds. Assuming the pseudorandomness of the hash function G , the probability of being $K \neq K'$ can be bounded by the DFP of SMAUG.PKE. The completeness of SMAUG.PKE (Theorem 2) concludes the proof. \square

4.3 Security proof

When proving the security of the KEMs constructed using FO transform in the (Q)ROM, one typically relies on the generic reductions from one-wayness or IND-CPA security of the underlying PKE. In the ROM, SMAUG.KEM has a tight reduction from the IND-CPA security of the underlying PKE, SMAUG.PKE. However, like other lattice-based constructions, the underlying PKE has a chance of decryption failures, which makes the generic reduction unapplicable [51] or non-tight [13, 41, 43] in the QROM. Therefore, we prove the IND-CCA security of SMAUG.KEM based on the non-tight QROM reduction of [13] as explained in Section 2 by proving the IND-CPA security of SMAUG.PKE.

Theorem 4 (IND-CPA security of SMAUG.PKE). *Assuming pseudorandomness of the underlying sampling algorithms, the IND-CPA security of SMAUG.PKE can be tightly reduced to the decisional MLWE and MLWR problems. Specifically, for any IND-CPA-adversary \mathcal{A} of SMAUG.PKE, there exist adversaries \mathcal{B}_0 ,*

\mathcal{B}_1 , \mathcal{B}_2 , and \mathcal{B}_3 attacking the pseudorandomness of XOF, and the pseudorandomness of sampling algorithms, the hardness of MLWE, and the hardness of MLWR, respectively, such that,

$$\begin{aligned} \text{Adv}_{\text{SMAUG.PKE}}^{\text{IND-CPA}}(\mathcal{A}) &\leq \text{Adv}_{\text{XOF}}^{\text{PR}}(\mathcal{B}_0) + \text{Adv}_{\text{expandA,HWT,dGaussian}}^{\text{PR}}(\mathcal{B}_1) \\ &\quad + \text{Adv}_{n,q,k,k}^{\text{MLWE}}(\mathcal{B}_2) + \text{Adv}_{n,p,q,k+1,k}^{\text{MLWR}}(\mathcal{B}_3). \end{aligned}$$

The secret distribution terms omitted in the last two advantages (of \mathcal{B}_1 and \mathcal{B}_2) are uniform over ternary polynomials with Hamming weights h_s and h_r , respectively. The error distribution term omitted in the advantage of \mathcal{B}_2 is a pseudorandom distribution following the corresponding CDT.

Proof. The proof proceeds by a sequence of hybrid games from G_0 to G_4 defined as follows:

- G_0 : the genuine IND-CPA game,
- G_1 : identical to G_0 , except that the public key is changed into (\mathbf{A}, \mathbf{b}) ,
- G_2 : identical to G_1 , except that the sampling algorithms are changed into truly random samplings,
- G_3 : identical to G_2 , except that \mathbf{b} is randomly chosen from \mathcal{R}_q^k ,
- G_4 : identical to G_3 , except that the ciphertext is randomly chosen from $\mathcal{R}_p^k \times \mathcal{R}_{p'}$. As a result, the public key and the ciphertexts are truly random.

We denote the advantage of the adversary on each game G_i as Adv_i , where $\text{Adv}_0 = \text{Adv}_{\text{SMAUG.PKE}}^{\text{IND-CPA}}(\mathcal{A})$ and $\text{Adv}_4 = 0$. Then, it holds that

$$|\text{Adv}_0 - \text{Adv}_1| \leq \text{Adv}_{\text{XOF}}^{\text{PR}}(\mathcal{B}_0),$$

for some adversary \mathcal{B}_0 against the pseudorandomness of the extendable output function. Given that the only difference between the transcripts viewed in hybrid games G_1 and G_2 is the randomness sampling, it can be concluded that

$$|\text{Adv}_1 - \text{Adv}_2| \leq \text{Adv}_{\text{expandA,HWT,dGaussian}}^{\text{PR}}(\mathcal{B}_1),$$

for some adversary, \mathcal{B}_1 attacking the pseudorandomness of at least one of the samplers. The difference in the games G_2 and G_3 is in the way the polynomial vector \mathbf{b} is sampled. In G_2 , it is sampled as part of an MLWE sample, whereas in G_3 , it is randomly selected. Thus, the difference in the advantages Adv_2 and Adv_3 can be bounded by $\text{Adv}_{n,q,k,k}^{\text{MLWE}}(\mathcal{B}_2)$, where \mathcal{B}_2 is an adversary distinguishing the MLWE samples from random. In the hybrids G_3 and G_4 , the only difference is in the way the ciphertexts are generated; they are either randomly chosen from $\mathcal{R}_p^k \times \mathcal{R}_{p'}$ or generated to be $(\mathbf{c}_1, \lfloor p'/p \cdot c_2 \rfloor)$, where

$$\begin{bmatrix} \mathbf{c}_1 \\ c_2 \end{bmatrix} = \left\lfloor \frac{p}{q} \cdot \begin{pmatrix} \mathbf{A} \\ \mathbf{b}^\top \end{pmatrix} \cdot \mathbf{r} \right\rfloor + \frac{p}{t} \cdot \begin{bmatrix} 0 \\ \mu \end{bmatrix}.$$

If an adversary \mathcal{A} can distinguish the two ciphertexts, we can construct an adversary \mathcal{B}_3 distinguishing the MLWR sample from random: *for given a sample*

$(\mathbf{A}, \mathbf{b}) \in \mathcal{R}_q^{(k+1) \times k} \times \mathcal{R}_p^{k+1}$, \mathcal{B}_3 rewrites \mathbf{b} as $(\mathbf{b}_1, b_2) \in \mathcal{R}_p^k \times \mathcal{R}_p$, computes $(\mathbf{b}_1, \lfloor p'/p \cdot b_2 \rfloor)$, and use \mathcal{A} to decide the ciphertext type. The output of \mathcal{A} will be the output of \mathcal{B}_3 . Therefore, we can conclude the proof by observing that

$$|\text{Adv}_3 - \text{Adv}_4| \leq \text{Adv}_{n,p,q,k+1,k}^{\text{MLWR}}(\mathcal{B}_3).$$

□

The classical IND-CCA security of SMAUG.KEM is then obtained directly from FO transforms [41] in the classical random oracle model. Theorem 1 implies the quantum IND-CCA security of SMAUG.KEM in the quantum random oracle model.

5 Parameter selection and concrete security

In this section, we first give a concrete security analysis of SMAUG and provide the recommended parameter sets.

5.1 Concrete security estimation

We exploit the best-known lattice attacks to estimate the concrete security of SMAUG.

5.1.1 Core-SVP methodology. Most of the known attacks are essentially finding a nonzero short vector in Euclidean lattices, using the Block–Korkine–Zolotarev (BKZ) lattice reduction algorithm [22, 40, 52]. BKZ has been used in various lattice-based schemes [3, 16, 31, 34, 55]. The security of the schemes is generally estimated as the time complexity of BKZ in core-SVP hardness introduced in [5]. It depends on the block size β of BKZ reporting the best performance. According to Becker et al. [10] and Chailloux et al. [21], the β -BKZ algorithm takes approximately $2^{0.292\beta+o(\beta)}$ and $2^{0.257\beta+o(\beta)}$ time in the classical and quantum setting, respectively. The polynomial factors and $o(\beta)$ terms in the exponent are ignored. We use the lattice estimator [2] to estimate the concrete security of SMAUG in core-SVP hardness.

5.1.2 Beyond Core-SVP methodology. In addition to lattice reduction attacks, we also take into consideration the cost of other types of attacks, e.g., algebraic attacks like the Arora-Ge attack or Coded-BKW attacks, and their variants. In general, these attacks have considerably higher costs and memory requirements compared to previously introduced attacks.

We also focus on the attacks not considered in the lattice estimator, specifically those that target sparse secrets, such as Meet-LWE [50] attack. This attack is inspired by Odlyzko’s Meet-in-the-Middle approach and involves using representations of ternary secrets in additive shares. The asymptotic attack complexity is claimed as $\mathcal{S}^{0.25}$; however, it is far from the estimated attack costs in SMAUG parameter sets. Even, the estimated cost has a significant gap with the real attack, due to the hidden costs behind the estimation.

We summarize the costs of the algebraic and combinatorial attacks in Table 2. Attack costs for Arora-Ge and Coded-BKW are estimated with lattice estimator [2]. The estimated cost of Arora-Ge attack on SMAUG 256 is not determined by lattice-estimator, outputting ∞ , which is at least a thousand bits of security. The costs for the Meet-LWE attack are estimated with a python script⁶ based on May’s analysis [50], best among Rep-1 and Rep-2.

⁶ The script can be found on the team SMAUG website: <http://kpqc.cryptolab.co.kr/>

Parameters sets		SMAUG 128	SMAUG 192	SMAUG 256
Target security		1	3	5
Classical core-SVP		120.0	181.7	264.5
Algebraic & Combinatorial attacks				
Arora-Ge	time	741.3	983.4	-
	(mem)	(598.0)	(636.5)	-
BKW	time	144.7	202.0	274.6
	(mem)	(133.7)	(190.7)	(256.9)
Meet-LWE	time	164.3	213.8	283.2
	(mem)	(143.7)	(192.4)	(254.7)

Table 2: Attack costs beyond Core-SVP.

5.1.3 MLWE hardness. We estimated the cost of the best-known attacks for MLWE, including *primal attack*, *dual attack*, and their hybrid variations, in the core-SVP hardness. We remark that any $\text{MLWE}_{n,q,k,\ell,\eta}$ instance can be viewed as an $\text{LWE}_{q,nk,n\ell,\eta}$ instance. Although the MLWE problem has an additional algebraic structure compared to the LWE problem, no attacks currently take advantage of this structure. Therefore, we assess the hardness of the MLWE problem based on the hardness of the corresponding LWE problem. We also consider the distributions of secret and noise when estimating the concrete security of SMAUG. We have also analyzed the costs of recent attacks that aim to target the MLWE problem with sparse secrets. Our narrow discrete Gaussian sampler’s tail bound is considered in estimating the security using the lattice estimator.

5.1.4 MLWR hardness. To measure the hardness of the MLWR problem, we treat it as an MLWE problem since no known attack utilizes the deterministic error term in the MLWR structure. Banerjee et al. [9] provided the reduction from the MLWE problem to the MLWR problem, which was subsequently improved in [6, 7, 15]. Basically, for given an MLWR sample $(\mathbf{A}, \lfloor p/q \cdot \mathbf{A} \cdot \mathbf{s} \rfloor \bmod p)$ with uniformly chosen $\mathbf{A} \leftarrow \mathcal{R}_q^k$ and $\mathbf{s} \leftarrow \mathcal{R}_p^\ell$, it can be expressed as $(\mathbf{A}, p/q \cdot (\mathbf{A} \cdot \mathbf{s} \bmod q) + \mathbf{e} \bmod p)$. The MLWR sample can be converted to an MLWE sample over \mathcal{R}_q by multiplying q/p as $(\mathbf{A}, \mathbf{b} = \mathbf{A} \cdot \mathbf{s} + q/p \cdot \mathbf{e} \bmod q)$. Assuming that the error term in the resulting MLWE sample is a random variable, uniformly distributed within the interval $(-q/2p, q/2p]$, we can estimate the hardness of the MLWR problem as the hardness of the corresponding MLWE problem.

5.2 Parameter sets

The SMAUG is parameterized by various integers such as n, k, q, p, p', t, h_s and h_r , as well as a standard deviation $\sigma > 0$ for the discrete Gaussian noise. Our main focus when selecting these parameters is to minimize the ciphertext size while maintaining security. We set SMAUG parameters to make SMAUG at least as safe as Saber. We first set our ring dimension to $n = 256$ and plaintext modulus to $t = 2$ to have a 256-bit message space (or sharing key space). Then we search for

parameters that offer the smallest ciphertext size with enough security. Starting from parameters having a tiny ciphertext size, we increase the ciphertext size, h_s , h_r , and σ and search for the parameters having enough security. Once we have them, we compute the DFP. If it is enough low, we can choose the compression parameter p' , but if it is not, we continue searching for appropriate parameters. The compression factor p' can be set to a small integer if the DFP is low enough. Else, we can keep $p' = 256$ as in the level-3 parameter, and not compress the ciphertext.

Table 3 shows the three parameter sets of **SMAUG**, corresponding to NIST’s security levels 1, 3, and 5. For security levels 3 and 5, we can not find the parameters for $q = 1024$, so we use $q = 2048$. Especially, the standard deviation $\sigma = 1.0625$ is too low for security level 3, so we move to $\sigma = 1.453713$. For the level-5 parameters set, we use $k = 5$ since $k = 4$ is too small for enough security.

Parameters sets Security level	SMAUG-128 1	SMAUG-192 3	SMAUG-256 5
n	256	256	256
k	2	3	5
(q, p, p', t)	(1024, 256, 32, 2)	(2048, 256, 256, 2)	(2048, 256, 64, 2)
(h_s, h_r)	(140, 132)	(198, 151)	(176, 160)
σ	1.0625	1.453713	1.0625
Classical core-SVP	120.0	181.7	264.5
Quantum core-SVP	105.6	160.9	245.2
Beyond core-SVP	144.7	202.0	274.6
DFP	-119.6	-136.1	-167.2
Secret key	176	236	218
Public key	672	1088	1792
Ciphertext	672	1024	1472

Table 3: The NIST security level, selected parameters, classical and quantum core-SVP hardness and security beyond core-SVP (see Section 5.1.2), DFP (in \log_2), and sizes (in bytes) of **SMAUG**.

The core-SVP hardness is estimated via the lattice estimator [2] using the cost model “ADPS16” introduced in [5] and “MATZOV” [49]. In the table, the smaller cost is reported. We assumed that the number of 1s is equal to the number of -1 s for simplicity, which conservatively underestimates security.

The security beyond core-SVP is estimated via the lattice estimator [2] and the Python script implementing the Meet-LWE attack cost estimation. It shows the lowest attack costs among coded-BKW, Arora-Ge, and Meet-LWE attack and their variants. We note that these attacks require a minimum memory size of 2^{130} to 2^{260} .

5.3 Decryption failure probability

As our primary goal is to push the efficiency of the lattice-based KEMs toward the limit while keeping roughly the same level of security, we follow the frameworks given in the NIST finalist Saber. In particular, we set the DFP to have similar to or lower than that of Saber’s.

The impact of DFP on the security of KEM is still being investigated. However, we can justify our decision to follow Saber’s choice and why it is sufficient for real-world scenarios. To do this, we make the following assumptions:

1. Each key pair has a limit of $Q_{\text{limit}} = 2^{64}$ decryption queries, as specified in NIST’s proposal call.
2. There are approximately 2^{33} people worldwide, each with hundreds of devices. Each device has hundreds of *usable* public keys broadcasted for KEM.
3. We introduce an observable probability and assume it is far less than 2^{-20} . Even though the decryption failure occurs, it can only be used for an attack when observed. Attackers can observe it through a side-channel attack, which enables the observation of decapsulation failures in the mounted device, or through direct communications after key derivation, allowing the detection of decryption failures with a communication per key pair. We assume the two cases can occur much less than 2^{-20} , as they require physically mounted devices or communications with shared keys.

Based on these assumptions, we can deduce that the number of observable decryption failures can be upper bounded by $2^{64+33+8+8} \cdot 2^{-20} = 2^{93}$. Based on the best-known (multi-target) attacks for Saber [27, Figure 6 (a)], the quantum cost for finding a single failing ciphertext of SMAUG security level 3 is much higher than 2^{160} , as desired⁷. For security level 5, we refer to Figure 7(a) in [27], which shows that the quantum cost for finding a single failure is much higher than 2^{245} . Regardless of the attack cost estimated above, the scenario of checking the failures in more than 2^{40} different devices is already way too far from the real-world attack scenario.

⁷ Specifically, the number of observable failures must be larger than $1/\beta$ in [27] to observe at least one failing ciphertext. That is, β should be larger than 2^{93} . The quantum cost is given as $1/\beta\sqrt{\alpha}$.

6 Implementation

In this section, we give the implementation performance for each parameter set. We compare the sizes and the reported performance with prior works such as Kyber and Saber. The constant-time reference implementation of SMAUG, along with the supporting scripts, can be found in our website: www.kpqc.cryptolab.co.kr/smaug, github: [www.kpqc.cryptolab.co.kr/smaug](https://github.com/kpqc/cryptolab-co.kr/smaug).

6.1 Performance

We instantiate the hash functions G, H and the extendable output function XOF with the following symmetric primitives: G is instantiated with SHAKE256, H is instantiated with SHA3-256, and XOF is instantiated with SHAKE128.

Table 4 presents the performance results of SMAUG. For a fair comparison, we also performed measurements on the same system with identical settings of the reference implementation of Kyber and Saber⁸. All benchmarks are obtained on one core of an Intel(R) Core(TM) i7-10700K CPU processor with clock speed 3.80GHz. The benchmarking machine has 64 GB of RAM and runs Debian GNU/Linux with Linux kernel version 5.4.0. The implementation is compiled with gcc version 9.4.0, and the compiler flags as indicated in the CMakeLists included in the submission package.

Schemes	Cycles			Cycles (ratio)		
	KeyGen	Encap	Decap	KeyGen	Encap	Decap
Kyber512	131560	162472	18930	1.8	2.27	2.08
LightSaber	93752	122176	133764	1.28	1.70	1.47
SMAUG-128	73061	71669	91007	1	1	1
Kyber768	214160	251308	285378	1.58	2.03	1.83
Saber	18722	224686	239590	1.38	1.82	1.53
SMAUG-192	135899	123763	156293	1	1	1
Kyber1024	332470	371854	415498	1.38	1.61	1.54
FireSaber	289278	347900	382326	1.20	1.51	1.41
SMAUG-256	240254	230521	270437	1	1	1

Table 4: Median cycle counts of 1000 executions for Kyber, Saber, and SMAUG (and their ratios). Cycle counts are obtained on one core of an Intel Core i7-10700k, with TurboBoost and hyperthreading disabled, using the C implementations without AVX optimizations.

6.2 Security against physical attacks

We justify the security of SMAUG against physical attacks based on the profiled Differential Power Analysis (DPA). Specifically, Simple Power Analysis (SPA)

⁸ From github.com/pq-crystals/kyber (518de24) and github.com/KULeuven-COSI C/SABER (f7f39e4), respectively.

can profile the key generation and encapsulation processes since they only occur once or without a secret key. However, multi-trace attacks are possible for decapsulation due to “re-encryption.” As Kyber and Saber share many design aspects with **SMAUG**, we can follow recent works on masking Kyber and Saber [11, 18] to add SCA countermeasures. Our new sampler, **dGaussian**, is expressed with bit operations, so adding SCA countermeasures like boolean masking is easy. While Krausz et al. [46] have recently proposed masking methods for the fixed hamming weight sampler, their efficiency is lacking, so we see it as future work. The new multiplication method may be vulnerable to memory access patterns, but we can efficiently mask it using coefficient-wise shuffling.

Acknowledgments Part of this work was done while MinJune Yi was in CryptoLab Inc. Part of this specification is from its conference version [23].

References

1. Akleylek, S., Alkim, E., Tok, Z.Y.: Sparse polynomial multiplication for lattice-based cryptography with small complexity. *The Journal of Supercomputing* **72**, 438–450 (2016)
2. Albrecht, M.R., Player, R., Scott, S.: On the concrete hardness of learning with errors. *Journal of Mathematical Cryptology* **9**(3), 169–203 (2015)
3. Alkim, E., Barreto, P.S.L.M., Bindel, N., Kramer, J., Longa, P., Ricardini, J.E.: The lattice-based digital signature scheme qtesla. *Cryptology ePrint Archive*, Paper 2019/085 (2019), <https://eprint.iacr.org/2019/085>
4. Alkim, E., Bos, J., Ducas, L., Longa, P., Mironov, I., Naehrig, M., Nikolaenko, V., Peikert, C., Raghunathan, A., Stebila, D.: Frodokem: Algorithm specifications and supporting documentation (2021)
5. Alkim, E., Ducas, L., Pöppelmann, T., Schwabe, P.: Post-quantum key exchange - A new hope. In: Holz, T., Savage, S. (eds.) *USENIX Security 2016*. pp. 327–343. *USENIX Association* (Aug 2016)
6. Alperin-Sheriff, J., Apon, D.: Dimension-preserving reductions from lwe to lwr. *Cryptology ePrint Archive*, Paper 2016/589 (2016), <https://eprint.iacr.org/2016/589>
7. Alwen, J., Krenn, S., Pietrzak, K., Wichs, D.: Learning with rounding, revisited. In: Canetti, R., Garay, J.A. (eds.) *Advances in Cryptology – CRYPTO 2013*. pp. 57–74. *Springer Berlin Heidelberg*, Berlin, Heidelberg (2013)
8. Arora, S., Ge, R.: New algorithms for learning in presence of errors. In: Aceto, L., Henzinger, M., Sgall, J. (eds.) *Automata, Languages and Programming*. pp. 403–415. *Springer Berlin Heidelberg*, Berlin, Heidelberg (2011)
9. Banerjee, A., Peikert, C., Rosen, A.: Pseudorandom functions and lattices. In: Pointcheval, D., Johansson, T. (eds.) *Advances in Cryptology – EUROCRYPT 2012*. pp. 719–737. *Springer Berlin Heidelberg*, Berlin, Heidelberg (2012)
10. Becker, A., Ducas, L., Gama, N., Laarhoven, T.: New directions in nearest neighbor searching with applications to lattice sieving, pp. 10–24. *Society for Industrial and Applied Mathematics* (2016). <https://doi.org/10.1137/1.9781611974331.ch2>
11. Beirendonck, M.V., D’anvers, J.P., Karmakar, A., Balasch, J., Verbauwhede, I.: A side-channel-resistant implementation of saber. *J. Emerg. Technol. Comput. Syst.* **17**(2) (apr 2021). <https://doi.org/10.1145/3429983>
12. Bi, L., Lu, X., Luo, J., Wang, K.: Hybrid dual and meet-LWE attack. In: Nguyen, K., Yang, G., Guo, F., Susilo, W. (eds.) *ACISP 22*. LNCS, vol. 13494, pp. 168–188. *Springer, Heidelberg* (Nov 2022). https://doi.org/10.1007/978-3-031-22301-3_9
13. Bindel, N., Hamburg, M., Hövelmanns, K., Hülsing, A., Persichetti, E.: Tighter proofs of CCA security in the quantum random oracle model. In: Hofheinz, D., Rosen, A. (eds.) *TCC 2019, Part II*. LNCS, vol. 11892, pp. 61–90. *Springer, Heidelberg* (Dec 2019). https://doi.org/10.1007/978-3-030-36033-7_3
14. Birkett, J., Dent, A.W.: Relations among notions of plaintext awareness. In: Cramer, R. (ed.) *Public Key Cryptography – PKC 2008*. pp. 47–64. *Springer Berlin Heidelberg*, Berlin, Heidelberg (2008)
15. Bogdanov, A., Guo, S., Masny, D., Richelson, S., Rosen, A.: On the hardness of learning with rounding over small modulus. In: Kushilevitz, E., Malkin, T. (eds.) *Theory of Cryptography*. pp. 209–224. *Springer Berlin Heidelberg*, Berlin, Heidelberg (2016)

16. Bos, J., Ducas, L., Kiltz, E., Lepoint, T., Lyubashevsky, V., Schanck, J.M., Schwabe, P., Seiler, G., Stehlé, D.: Crystals-kyber: a cca-secure module-lattice-based kem. In: 2018 IEEE European Symposium on Security and Privacy (EuroS&P). pp. 353–367. IEEE (2018)
17. Bos, J.W., Costello, C., Ducas, L., Mironov, I., Naehrig, M., Nikolaenko, V., Raghunathan, A., Stebila, D.: Frodo: Take off the ring! Practical, quantum-secure key exchange from LWE. In: Weippl, E.R., Katzenbeisser, S., Kruegel, C., Myers, A.C., Halevi, S. (eds.) ACM CCS 2016. pp. 1006–1018. ACM Press (Oct 2016). <https://doi.org/10.1145/2976749.2978425>
18. Bos, J.W., Gourjon, M., Renes, J., Schneider, T., van Vredendaal, C.: Masking kyber: First- and higher-order implementations. IACR TCHES **2021**(4), 173–214 (2021). <https://doi.org/10.46586/tches.v2021.i4.173-214>, <https://tches.iacr.org/index.php/TCHES/article/view/9064>
19. Bossuat, J.P., Troncoso-Pastoriza, J.R., Hubaux, J.P.: Bootstrapping for approximate homomorphic encryption with negligible failure-probability by using sparse-secret encapsulation. In: Ateniese, G., Venturi, D. (eds.) ACNS 22. LNCS, vol. 13269, pp. 521–541. Springer, Heidelberg (Jun 2022). https://doi.org/10.1007/978-3-031-09234-3_26
20. Boudgoust, K., Jeudy, C., Roux-Langlois, A., Wen, W.: On the hardness of module learning with errors with short distributions. Journal of Cryptology **36**(1), 1 (Jan 2023). <https://doi.org/10.1007/s00145-022-09441-3>
21. Chailloux, A., Loyer, J.: Lattice sieving via quantum random walks. In: Tibouchi, M., Wang, H. (eds.) Advances in Cryptology - ASIACRYPT. pp. 63–91 (2021)
22. Chen, Y., Nguyen, P.Q.: BKZ 2.0: Better lattice security estimates. In: Lee, D.H., Wang, X. (eds.) ASIACRYPT 2011. LNCS, vol. 7073, pp. 1–20. Springer, Heidelberg (Dec 2011). https://doi.org/10.1007/978-3-642-25385-0_1
23. Cheon, J.H., Choe, H., Hong, D., Yi, M.: Smaug: Pushing lattice-based key encapsulation mechanisms to the limits. Cryptology ePrint Archive, Paper 2023/739 (2023), <https://eprint.iacr.org/2023/739>
24. Cheon, J.H., Han, K., Kim, A., Kim, M., Song, Y.: Bootstrapping for approximate homomorphic encryption. In: Nielsen, J.B., Rijmen, V. (eds.) EUROCRYPT 2018, Part I. LNCS, vol. 10820, pp. 360–384. Springer, Heidelberg (Apr / May 2018). https://doi.org/10.1007/978-3-319-78381-9_14
25. Cheon, J.H., Han, K., Kim, J., Lee, C., Son, Y.: A practical post-quantum public-key cryptosystem based on splWE. In: Hong, S., Park, J.H. (eds.) ICISC 16. LNCS, vol. 10157, pp. 51–74. Springer, Heidelberg (Nov / Dec 2017). https://doi.org/10.1007/978-3-319-53177-9_3
26. Cheon, J.H., Kim, D., Lee, J., Song, Y.: Lizard: Cut off the tail! A practical post-quantum public-key encryption from LWE and LWR. In: Catalano, D., De Prisco, R. (eds.) SCN 18. LNCS, vol. 11035, pp. 160–177. Springer, Heidelberg (Sep 2018). https://doi.org/10.1007/978-3-319-98113-0_9
27. D’Anvers, J.P., Batsleer, S.: Multitarget decryption failure attacks and their application to saber and kyber. In: Hanaoka, G., Shikata, J., Watanabe, Y. (eds.) PKC 2022, Part I. LNCS, vol. 13177, pp. 3–33. Springer, Heidelberg (Mar 2022). https://doi.org/10.1007/978-3-030-97121-2_1
28. D’Anvers, J.P., Guo, Q., Johansson, T., Nilsson, A., Vercauteren, F., Verbauwhede, I.: Decryption failure attacks on ind-cca secure lattice-based schemes. In: Lin, D., Sako, K. (eds.) Public-Key Cryptography – PKC 2019. pp. 565–598. Springer International Publishing, Cham (2019)

29. D’Anvers, J.P., Karmakar, A., Roy, S.S., Vercauteren, F.: Saber: Module-LWR based key exchange, CPA-secure encryption and CCA-secure KEM. In: Joux, A., Nitaj, A., Rachidi, T. (eds.) AFRICACRYPT 18. LNCS, vol. 10831, pp. 282–305. Springer, Heidelberg (May 2018). https://doi.org/10.1007/978-3-319-89339-6_16
30. Dent, A.W.: A designer’s guide to kems. In: Cryptography and Coding: 9th IMA International Conference, Cirencester, UK, December 16-18, 2003. Proceedings 9. pp. 133–151. Springer (2003)
31. Ducas, L., Kiltz, E., Lepoint, T., Lyubashevsky, V., Schwabe, P., Seiler, G., Stehlé, D.: CRYSTALS-Dilithium: A lattice-based digital signature scheme. IACR TCHES **2018**(1), 238–268 (2018). <https://doi.org/10.13154/tches.v2018.i1.238-268>, <https://tches.iacr.org/index.php/TCHES/article/view/839>
32. D’Anvers, J.P., Karmakar, A., Sinha Roy, S., Vercauteren, F.: Saber: Module-lwr based key exchange, cpa-secure encryption and cca-secure kem. In: International Conference on Cryptology in Africa. pp. 282–305. Springer (2018)
33. Espitau, T., Joux, A., Kharchenko, N.: On a dual/hybrid approach to small secret LWE - A dual/enumeration technique for learning with errors and application to security estimates of FHE schemes. In: Bhargavan, K., Oswald, E., Prabhakaran, M. (eds.) INDOCRYPT 2020. LNCS, vol. 12578, pp. 440–462. Springer, Heidelberg (Dec 2020). https://doi.org/10.1007/978-3-030-65277-7_20
34. Fouque, P.A., Hoffstein, J., Kirchner, P., Lyubashevsky, V., Pornin, T., Prest, T., Ricosset, T., Seiler, G., Whyte, W., Zhang, Z.: Falcon: Fast-fourier lattice-based compact signatures over ntru. Submission to the NIST’s post-quantum cryptography standardization process **36**(5) (2018)
35. Fujisaki, E., Okamoto, T.: Secure integration of asymmetric and symmetric encryption schemes. In: Wiener, M.J. (ed.) CRYPTO’99. LNCS, vol. 1666, pp. 537–554. Springer, Heidelberg (Aug 1999). https://doi.org/10.1007/3-540-48405-1_34
36. Fujisaki, E., Okamoto, T.: Secure integration of asymmetric and symmetric encryption schemes. Journal of Cryptology **26**(1), 80–101 (Jan 2013). <https://doi.org/10.1007/s00145-011-9114-1>
37. Grover, L.K.: A fast quantum mechanical algorithm for database search. In: Proceedings of the twenty-eighth annual ACM symposium on Theory of computing. pp. 212–219 (1996)
38. Guo, Q., Johansson, T., Stankovski, P.: Coded-bkw: Solving lwe using lattice codes. In: Annual Cryptology Conference. pp. 23–42. Springer (2015)
39. Halevi, S., Shoup, V.: Bootstrapping for HELib. In: Oswald, E., Fischlin, M. (eds.) EUROCRYPT 2015, Part I. LNCS, vol. 9056, pp. 641–670. Springer, Heidelberg (Apr 2015). https://doi.org/10.1007/978-3-662-46800-5_25
40. Hanrot, G., Pujol, X., Stehlé, D.: Algorithms for the shortest and closest lattice vector problems. In: Chee, Y.M., Guo, Z., Ling, S., Shao, F., Tang, Y., Wang, H., Xing, C. (eds.) Coding and Cryptology. pp. 159–190. Springer Berlin Heidelberg, Berlin, Heidelberg (2011)
41. Hofheinz, D., Hövelmanns, K., Kiltz, E.: A modular analysis of the Fujisaki-Okamoto transformation. In: Kalai, Y., Reyzin, L. (eds.) TCC 2017, Part I. LNCS, vol. 10677, pp. 341–371. Springer, Heidelberg (Nov 2017). https://doi.org/10.1007/978-3-319-70500-2_12
42. Hong, S.: Lizarderror. <https://github.com/swanhong/LizardError> (2018)
43. Hövelmanns, K., Kiltz, E., Schäge, S., Unruh, D.: Generic authenticated key exchange in the quantum random oracle model. In: Kiayias, A., Kohlweiss, M., Wallden, P., Zikas, V. (eds.) PKC 2020, Part II. LNCS, vol. 12111, pp. 389–422.

- Springer, Heidelberg (May 2020). https://doi.org/10.1007/978-3-030-45388-6_14
44. Howgrave-Graham, N., Nguyen, P.Q., Pointcheval, D., Proos, J., Silverman, J.H., Singer, A., Whyte, W.: The impact of decryption failures on the security of ntru encryption. In: Boneh, D. (ed.) *Advances in Cryptology - CRYPTO 2003*. pp. 226–246. Springer Berlin Heidelberg, Berlin, Heidelberg (2003)
 45. Jiang, H., Zhang, Z., Chen, L., Wang, H., Ma, Z.: Ind-cca-secure key encapsulation mechanism in the quantum random oracle model, revisited. In: Shacham, H., Boldyreva, A. (eds.) *Advances in Cryptology – CRYPTO 2018*. pp. 96–125. Springer International Publishing, Cham (2018)
 46. Krausz, M., Land, G., Richter-Brockmann, J., Güneysu, T.: A holistic approach towards side-channel secure fixed-weight polynomial sampling. In: *Public-Key Cryptography–PKC 2023: 26th IACR International Conference on Practice and Theory of Public-Key Cryptography*, Atlanta, GA, USA, May 7–10, 2023, *Proceedings, Part II*. pp. 94–124. Springer (2023)
 47. Langlois, A., Stehlé, D., Steinfeld, R.: GGHLite: More efficient multilinear maps from ideal lattices. In: Nguyen, P.Q., Oswald, E. (eds.) *EUROCRYPT 2014*. LNCS, vol. 8441, pp. 239–256. Springer, Heidelberg (May 2014). https://doi.org/10.1007/978-3-642-55220-5_14
 48. Lee, J., Kim, D., Lee, H., Lee, Y., Cheon, J.H.: Rlizard: Post-quantum key encapsulation mechanism for iot devices. *IEEE Access* **7**, 2080–2091 (2018)
 49. MATZOV: Report on the Security of LWE: Improved Dual Lattice Attack (Apr 2022). <https://doi.org/10.5281/zenodo.6493704>
 50. May, A.: How to meet ternary LWE keys. In: Malkin, T., Peikert, C. (eds.) *CRYPTO 2021, Part II*. LNCS, vol. 12826, pp. 701–731. Springer, Heidelberg, Virtual Event (Aug 2021). https://doi.org/10.1007/978-3-030-84245-1_24
 51. Saito, T., Xagawa, K., Yamakawa, T.: Tightly-secure key-encapsulation mechanism in the quantum random oracle model. In: Nielsen, J.B., Rijmen, V. (eds.) *EUROCRYPT 2018, Part III*. LNCS, vol. 10822, pp. 520–551. Springer, Heidelberg (Apr / May 2018). https://doi.org/10.1007/978-3-319-78372-7_17
 52. Schnorr, C.P., Euchner, M.: Lattice basis reduction: Improved practical algorithms and solving subset sum problems. *Mathematical programming* **66**(1), 181–199 (1994)
 53. Sendrier, N.: Secure sampling of constant-weight words – application to bike. *Cryptology ePrint Archive*, Paper 2021/1631 (2021), <https://eprint.iacr.org/2021/1631>
 54. Son, Y., Cheon, J.H.: Revisiting the hybrid attack on sparse secret lwe and application to he parameters. In: *Proceedings of the 7th ACM Workshop on Encrypted Computing & Applied Homomorphic Cryptography*. p. 11–20. WAHC’19, Association for Computing Machinery, New York, NY, USA (2019). <https://doi.org/10.1145/3338469.3358941>
 55. Vercauteren, I.F., Sinha Roy, S., D’Anvers, J.P., Karmakar, A.: Saber: Mod-lwr based kem, nIST PQC Round 3 Submission